

VIETNAM NATIONAL UNIVERSITY, HO CHI MINH CITY
HO CHI MINH CITY UNIVERSITY OF SCIENCE
FACULTY OF INFORMATION TECHNOLOGY



STATISTICAL MACHINE LEARNING

FINAL PROJECT

OBJECT DETECTION WITH YOLO

Students: 21127004 - Tran Nguyen An Phong
21127135 - Diep Huu Phuc
21127307 - Tong Gia Huy
Class: 21CNTThuc
Instructors: Le Long Quoc
Ngo Minh Nhut

Ho Chi Minh City, August 2024

Abstract

Computer vision has long been one of the main foci in the field of computer science. This branch aims to utilize images or videos, and the computational power of machineries to perform tasks like classification, object detection, and recognition. Since the invention of convolutional neural network (CNN), computer vision models have been enhanced tremendously. Different from normal feed-forward networks, CNN utilizes kernels to extract regional features from the image, thus making calculations more accurate. Architectures like You Only Look Once (YOLO) leverages this advantage of preserving dimensionality to perform robust and high precision object detection. In this project, we aim to establish an overall understanding of this series of models, as well as use it to construct our own detection website. We will discuss the technicalities and notable features of YOLO, as well as give an overview of our website and details on our own model.

TABLE OF CONTENTS

1	INTRODUCTION	3
2	OVERVIEW ON OBJECT DETECTION MODELS	3
2.1	Architecture	3
2.2	Additional features	4
3	REQUIREMENT 1: Web interface for pre-built object detection model	4
4	REQUIREMENT 2: Custom model training and application building	4
4.1	Data Preparation	4
4.1.1	Data source	4
4.1.2	Data distribution	5
4.2	Training	6
4.2.1	Basics on training models with YOLOv4	6
4.2.2	Training YOLOv8 models with Ultralytics and Kaggle	7
4.3	Result and Evaluation	8
4.3.1	Evaluation of the training process	9
4.3.2	Evaluation of the testing process	12
4.4	Application	13
	ACKNOWLEDGMENT	16
	REFERENCES	16

1 INTRODUCTION

You Only Look Once or YOLO (Redmon et al. 2016) is a well-known series of object detection models. Throughout the years, more revised versions have been developed and published, with each later model having more modifications thus boosting its accuracy and performance.

In this project, we aim to establish an understanding of YOLO's architecture. We then used our own dataset to retrain the model, and developed a web interface for object detection. In this report, we present our overview on YOLO's features and our website's development. Details on our training process, including data, evaluation metrics, and results are also discussed. Before going in-depth, here is a table summarizing everything we have achieved.

Table 1. Task assignment table

ID	Name	Task	Status (%)
21127004	Tran N. An Phong	<i>Requirement 2:</i> - Crawl data and compile datasets. - Implement YOLOv4. - Train and evaluate YOLOv8 models.	100
		<i>Report:</i> - Section 4.1 and 4.2.	
21127135	Diep Huu Phuc	<i>Requirement 1:</i> - Find model, and build Home page. <i>Requirement 2:</i> - Build Home page and Wiki page. - Evaluate YOLOv8 models.	100
		<i>Report:</i> - Section 3 and 4.4. - Proofread and finalize.	
21127307	Tong Gia Huy	<i>Requirement 2:</i> - Implement and train YOLOv4. - Design Wiki page. - Compile animal facts and links.	100
		<i>Report:</i> - Abstract, Introduction, Acknowledgement. - Section 2 and 4.3. - Proofread and finalize.	

2 OVERVIEW ON OBJECT DETECTION MODELS

2.1 Architecture

Object detection models have a number of components that help extract features, perform detection, and some additional supporting modules. In this section, we give an overview of the architecture of such models.

Backbone. As the name suggests, this part of the model is very important to the whole architecture as it is tasked with extracting the features from input images. Instead of encoding the features into a vector (like in convolutional networks) for classification, the backbone network saves the features into a stacked feature maps. This approach is crucial for preserving spacial information of the pixels, which allows for accurate detection of objects throughout the entire

image. Some widely used networks for this task include: CSPDarknet53 (Wang et al. 2020) used by Bochkovskiy et al. 2020 in YOLOv4, and EfficientNet (Tan 2019) used in YOLOv8.

Neck. This section is an optional part of the model and lies between the feature extraction stage (backbone) and detection stage (head). The purpose of this part is to construct multiple scales of the feature maps obtained from backbone. By expanding the dimension of the feature maps, the head can detect objects at different sizes. Some widely used methods for this part are: Spatial Pyramid Pooling (He et al. 2015) and Path Aggregation Network (Liu et al. 2018) in YOLOv4,

Head. One last main component of the object detection is responsible for performing the actual detection on the feature maps. For architectures that make use of the neck stage for enriching the scales, detection is performed at each scale independently from each other. Small scale maps are used for the detection of large objects, and larger scale maps are used for smaller objects'. It is in this stage that the bounding box for each object are defined, a classification probability is also calculated.

2.2 Additional features

In addition to the base architecture, some techniques for boosting the model's performance can also be incorporated during various stages. Bochkovskiy et al. 2020 call these adjustments *Bag of Freebies* and *Bag of Specials*. In this section, we give a simple explanation and example for each of those terms.

- Bag of Freebies: a set of techniques or methods that can alter the training strategy to improve the prediction accuracy, namely: Data Augmentation - used to alter the input images which can help create diversity in training data, Semantic Distribution Bias, and Objective Function of Bounding Box Regression.
- Bag of Specials: a set of modules and post-processing methods that can further improve the model's accuracy exchanging for the model's inference time. These methods can range from enhancing features to introducing advanced mechanisms.

3 REQUIREMENT 1: Web interface for pre-built object detection model

In this section, we present the design of our web interface, along with the integration of a pre-trained YOLO model to perform object detection. We used Ultralytics' version 10 medium model for the detection task. The model is loaded into global variables at the website's booting stage. If the pre-trained .pt file is not found on the local machine, it will be automatically downloaded. We present screenshots of the interface in Figure 1 and 2.

4 REQUIREMENT 2: Custom model training and application building

4.1 Data Preparation

In this part we will discuss about the data source and its distribution information. We will not delve deeply into the coding aspect as it is messy and also unnecessary for this report.

4.1.1 Data source

Although there are many image detection datasets available online (mainly on Kaggle), most of them do not meet our requirements which are the ability to filter by class (animal in our case) and

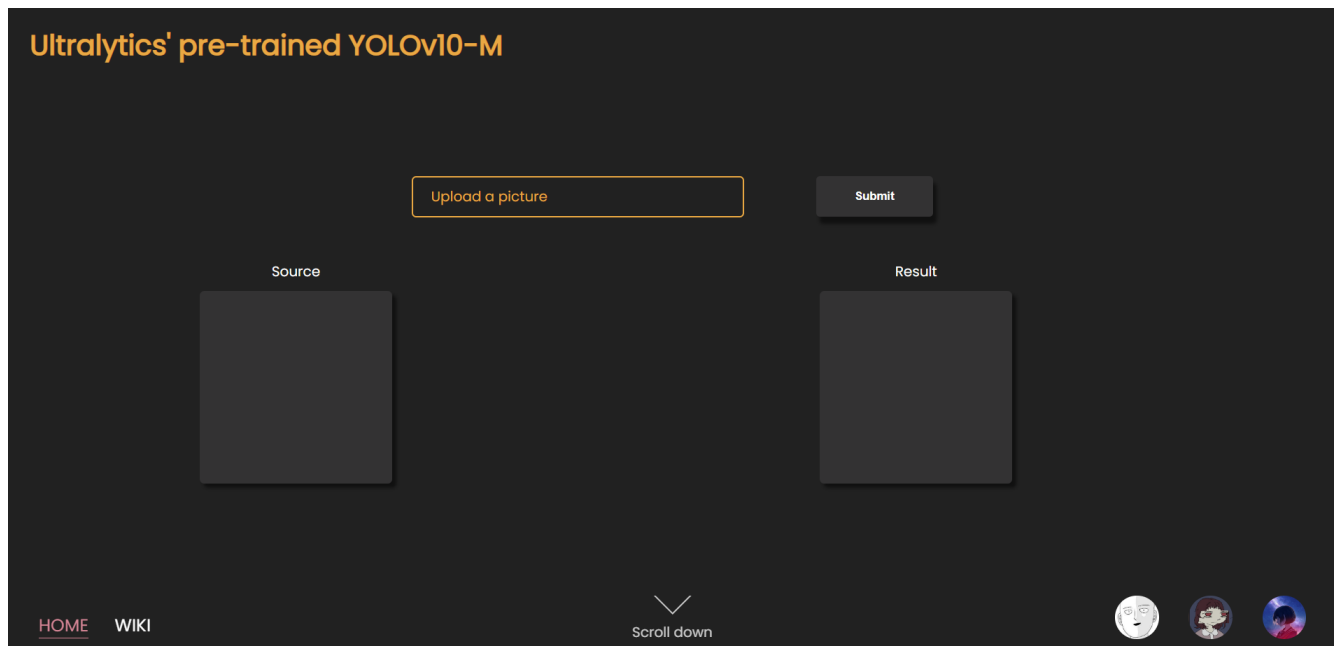


Figure 1. Default user interface for object detection task using Ultralytics’s pre-trained model.

having labels written in YOLO format. Therefore, we must come up with a solution to acquire the datasets that are as tailored to our needs as possible. This will be done with the help of [Google Open Images Dataset V7](#) and the [fiftyone](#) library.

YOLOv4: Implemented with Pytorch, it performs detection with box format as $[id, x1, y1, x2, y2]$ (this is different from both Google Open Images Dataset V7 and the YOLO format). Fortunately, there exist Kaggle datasets with this format, and as we don’t have high demands for this model, it is convenient to leverage resources that are already available. Hence, we will use a subset of the [Animals Detection Images Dataset](#) to work with this model.

YOLOv8: As mentioned above, our image database will be Google Open Images Dataset V7, and the required data will be crawled and transformed into YOLO format with the [fiftyone](#) library. This way, data can be extracted freely, with all required classes, and is only limited by what the database can provide. There is also another advantage of crawling our own data, which will be explained in the next part.

4.1.2 Data distribution

YOLOv4: A subset of 9 classes is taken from the original dataset, which has highly uneven distribution (Figure 4). This is one of the reasons why data needs to be crawled for YOLOv8. There is a total of 1858 images for training and testing.

YOLOv8: We choose to pick 22 classes from the database and divide it into three complete divisions of train/test/validation, which each have 7094/3691/1692 images respectively. So, the ratio for train/test/validation is nearly 57/30/13.

Figure 5 is a distribution plot for the training set, which is more evenly distributed thanks to the help of the crawling tool. Although there remain some classes that have a small number of images due to the providing limit of Google Open Image Dataset, they still possess reasonable quantities, which is enough to avoid the bias of providing too many or too little data for a class.




Figure 2. Interface for Ultralytics' model after being prompted with an input image.

4.2 Training

4.2.1 Basics on training models with YOLOv4

The YOLO architecture comprises of a Backbone network (to extract features from image) and a Detection head (to predict bounding boxes, objectness scores, and class probabilities), we first need to train the backbone network so the extracted information can prove more useful for the detection layer. Then, we apply regression training using a loss function for the Detection head so it can predict bounding boxes and performing classification better at the same time.

A training process goes through these steps:

1. Forward Pass: The input image is passed through the model. The model predicts bounding boxes, objectness scores, and class probabilities for each grid cell.
2. Loss Calculation:
 - Bounding Box Loss: Measures the error between the predicted bounding boxes and the ground truth bounding boxes using mean squared error.
 - Objectness Loss: Evaluates how well the model predicts the presence of an object in a bounding box using binary cross-entropy loss.
3. Class Loss: Measures the error between the predicted probabilities and the true class' using cross-entropy loss.
4. Backpropagation: The total loss is backpropagated through the network to update the weights using an optimization algorithm like Stochastic Gradient Descent (SGD) or Adam.
5. Iteration: This process is repeated for multiple epochs until the model converges to a set of weights that minimizes the loss function.

When talking about training a YOLO model, we often skip the Backbone training part as the Backbone network is highly universal and can be trained once for many different Detection

```
def download_animal(animal):  
    classes = []  
    classes.append(animal)  
    dataset = fiftyone.zoo.load_zoo_dataset(  
        "open-images-v7",  
        split="test",  
        label_types=["detections"],  
        classes=classes,  
        max_samples=400,  
        dataset_name="Animal_test_" + animal  
    )  
    dataset.export(  
        dataset_type=fiftyone.types.YOLOv5Dataset,  
        classes=classes,  
        split="test"  
    )
```

Figure 3. A snippet of code using `fiftyone` to download and convert data to YOLO format.

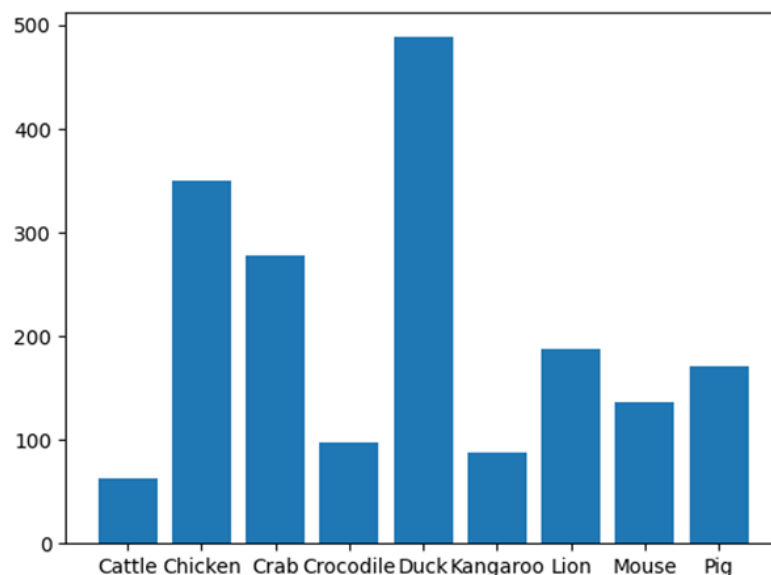


Figure 4. Distribution of the data used to train our YOLO4 model.

heads. With that, it is recommended to transfer the Backbone network's weights from notably well-trained models (e.g., Darknet53) instead of training them from scratch. Now, the training process mainly focuses on training the Detector.

YOLOv4: Because we only use this model for further insights of how a YOLO model training process works, it is only trained for 100 epochs. The training code is implemented based roughly on the process mentioned above, so it is not yet polished and many flaws may still remain. However, this should be acceptable considering the research nature and self-taught purpose of this model. The code used for this stage is referenced from [Tianxiaomo 2021](#).

4.2.2 Training YOLOv8 models with Ultralytics and Kaggle

From this point onward, we will employ the Ultralytics framework, and because of the demanding training process, everything is going to be done and hosted on Kaggle. A training process with Ultralytics has following steps:

1. Prepare datasets in YOLO format (which is discussed in Section 4.1).




Figure 5. Distribution of the data used to train our YOLO4 model.

2. Prepare a configuration file to guide Ultralytics for the new classes and train/val datasets.
3. Host these datasets and configuration file on Kaggle, where the training process will be carried out.
4. Tell Ultralytics to carry the Backbone network's weights from its base YOLOv8 model, and use our new configuration to train for the animal detection problem. While training, we need to specify the number of epochs and the input image size of the training dataset, which in our case are 200 epochs and 640.
5. When the training process is completed on Kaggle, the notebook returns outputs containing the trained model and several insight reports of the training process.

4.3 Result and Evaluation

For the sake of simplicity, we are going to evaluate only the YOLOv8 model trained with Ultralytics, as this is the model with the highest potential for a satisfactory evaluation score. Since we are working with an object detection model, the most suitable metric for grading would be mAP. Mean Average Precision (mAP) score is a popular evaluation metric, used in object detection models to measure their accuracy. It is particularly useful because both the precision and recall of the model's predictions are considered. Following is a breakdown of what the mAP score is and how it is calculated:

- The precision-recall curve (Figure 6) shows the trade-off between precision and recall for different threshold values. As you increase the threshold for considering a detection as positive, precision generally increases, but recall decreases, and vice versa.
- AP is the area under the precision-recall curve for a particular class. It is calculated by taking the average of precision values across all recall levels from 0 to 1. In some cases, interpolation is used to ensure that precision is monotonically decreasing. For each object class, the AP is calculated separately.
- mAP is the mean of the Average Precision (AP) values across all object classes. It provides a single score that summarizes the model's overall performance across all classes.
- Ultralytics definition of mAP is a little bit different, they define class' mAP as a mean between all "easy detections" of one class. There is also a model's mAP, which is the mean of all classes' mAP values. Ultralytics provides mAP score calculated under many IoU

(intersection over union) thresholds, which is a way of measuring how well the predicted bounding box overlaps with the ground truth. We choose mAP-50 as the common metric in this report.

4.3.1 Evaluation of the training process

Ultralytics provides mAP-50 score evaluated per epoch using the validation set. Since we are training the model on 200 epochs, it is redundant to show mAP score for every epoch in this report. Instead, we note down a mAP-50 score every 20 epochs in Table 2 (mAP-50 score for every epoch can still be accessed through the notebook output, or the “`result_per_epoch.csv`” file in the “`modelEvaluating`” folder).

Table 2. mAP-50s of the model while training. Measurements are taken every 20 epochs.

Epoch num	mAP-50
20	0.50381
40	0.65829
60	0.71419
80	0.736
100	0.74504
120	0.75555
140	0.76199
160	0.76561
180	0.76819
200	0.76527

As can be observed from Figure 6 and 7, all the metrics of our model (including mAP-50) saw significant changes while training less than 100 epochs. After 150 epochs, while there are still notable changes, we can almost say that the value hits its plateau. Maybe with the current dataset and configuration, we just need to train our model for 150 epochs.

In summary, our trained model achieves the mAP-50 score of 0.769, which is a fairly good point (a model with mAP-50 score above 0.5 can be deemed usable). Almost all classes have their own mAP-50 scores above 0.5, except for some special cases like Bear, Cattle or Sheep. This is possibly due to the low number of training samples and the vast number of animals that have similar features. For instance, for camels that have thick fur coats, the model can mistake them for sheep, and the same goes for other animals that have similarities like horns, skin colour, shape, etc. Therefore, the training data also plays an important role in determining the performance of the model. The resolution of the input image can also greatly affect the accuracy as unclear pictures can confuse the model by making it unable to confidently detect and identify the subject.




Figure 6. Precision-recall curve.




Figure 7. Metric report of the model while training.

Model summary (fused): 168 layers, 3,009,938 parameters, 0 gradients, 8.1 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95)	
all	1691	2260	0.769	0.711	0.769	0.614	100% 106/106 [00:15<00:00, 6.93it/s]
Bear	21	23	0.43	0.478	0.515	0.462	
Camel	25	30	0.735	0.833	0.856	0.747	
Cat	345	389	0.962	0.78	0.908	0.756	
Cattle	134	259	0.62	0.405	0.445	0.328	
Chicken	49	72	0.828	0.738	0.809	0.651	
Crab	66	72	0.916	0.875	0.931	0.697	
Crocodile	39	53	0.92	0.655	0.82	0.595	
Dog	400	480	0.92	0.624	0.845	0.705	
Duck	68	149	0.943	0.661	0.787	0.573	
Elephant	33	56	0.782	0.786	0.84	0.681	
Fox	41	44	0.529	0.795	0.718	0.671	
Frog	39	48	0.861	0.833	0.883	0.721	
Hippopotamus	10	13	0.535	0.846	0.746	0.664	
Kangaroo	21	32	0.591	0.719	0.666	0.52	
Lion	57	65	0.827	0.769	0.819	0.684	
Mouse	45	55	0.89	0.734	0.833	0.513	
Pig	51	76	0.757	0.592	0.669	0.539	
Rabbit	64	71	0.855	0.748	0.856	0.72	
Sheep	62	122	0.37	0.402	0.361	0.261	
Snail	34	40	0.917	0.9	0.941	0.744	
Snake	61	66	0.903	0.707	0.897	0.654	
Zebra	26	45	0.829	0.753	0.772	0.632	

Speed: 0.2ms preprocess, 1.9ms inference, 0.0ms loss, 1.0ms postprocess per image
Results saved to runs/detect/train

Figure 8. Metric report of the model after running on the validation set.

4.3.2 Evaluation of the testing process

In addition to training evaluation, we utilize a separate testing set to check our model's performance on new and unseen data. Ultralytics provides a `val()` method for valuation after training (Figure 9).

```
from ultralytics import YOLO

# Load a model
model = YOLO("/kaggle/input/yolov8-model-after-animal-training/last.pt")

# Customize validation settings
validation_results = model.val(data="/kaggle/input/yolov8-config-for-test/YOLOv8_for_Test.yaml", device="0")
```

Figure 9. Calling validation method with the trained model and test dataset.

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	231/231 [00:27<00:00, 8.31it/s]
all	3690	5188	0.733	0.685	0.715	0.565	
Bear	59	70	0.482	0.386	0.37	0.311	
Camel	56	91	0.611	0.681	0.603	0.489	
Cat	400	438	0.861	0.785	0.866	0.722	
Cattle	388	765	0.648	0.41	0.435	0.32	
Chicken	176	240	0.848	0.767	0.826	0.72	
Crab	157	195	0.874	0.759	0.787	0.603	
Crocodile	96	108	0.889	0.816	0.869	0.642	
Dog	400	489	0.887	0.624	0.827	0.687	
Duck	199	382	0.879	0.652	0.78	0.601	
Elephant	137	234	0.606	0.704	0.673	0.546	
Fox	103	110	0.732	0.873	0.835	0.742	
Frog	111	131	0.809	0.702	0.824	0.607	
Hippopotamus	39	58	0.548	0.741	0.593	0.449	
Kangaroo	65	87	0.571	0.701	0.679	0.545	
Lion	138	156	0.813	0.754	0.819	0.709	
Mouse	210	245	0.791	0.664	0.747	0.467	
Pig	135	212	0.688	0.655	0.704	0.607	
Rabbit	215	267	0.831	0.683	0.769	0.636	
Sheep	179	396	0.453	0.475	0.378	0.282	
Snail	129	139	0.721	0.755	0.77	0.603	
Snake	234	243	0.9	0.848	0.905	0.659	
Zebra	64	132	0.672	0.644	0.668	0.483	

Speed: 0.2ms preprocess, 2.2ms inference, 0.0ms loss, 1.1ms postprocess per image
Results saved to runs/detect/val

Figure 10. Evaluation metrics of our trained model on test dataset.

It is delightful to know that our trained model can achieve a mAP-50 score of 0.715 on the testing set, which is relatively close to the validation's discussed earlier. Considering the large size of the testing set, this result is acceptable. To a certain extent, our trained model is more than capable of accurately detecting animals on larger unseen dataset.

4.4 Application

Here we describe a practical usage of the trained model by incorporating it in building an animal detection website. In terms of technicalities, the model is loaded and used the same way as described in Section 3, with the exception of automatically downloading the trained weights.

We summarize the features of the site as follows:

- Perform wildlife animal detection on input image.
- List facts and related links on detected animals on the Wiki tab.
- Order the facts based on the confidence of each detection or based on the animal names.

To access the wildlife detection model, simply scroll down from the initial Ultralytics' homepage, or click on the *Scroll down* button on the bottom of the screen. To see the facts/links related to the animals, click on the *Wiki* button. Screenshots of the webpage are presented in Figure 11, 12, 13, and 14.




Figure 11. Interface of wildlife detection model's homepage.

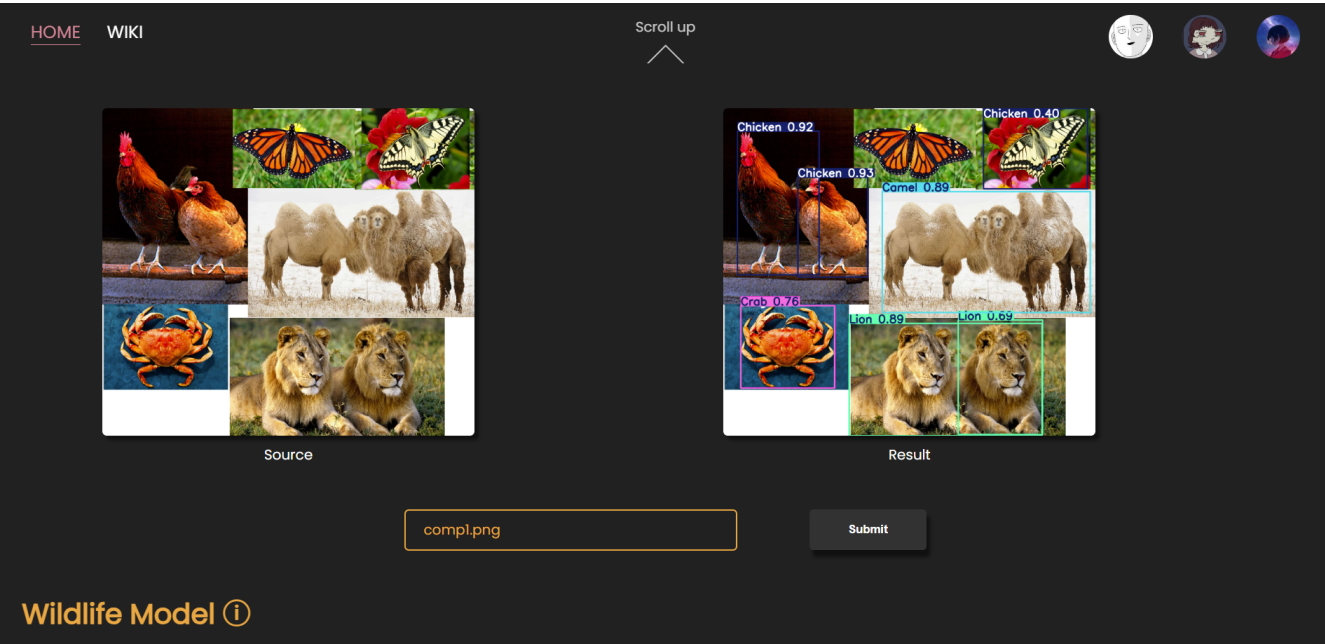


Figure 12. Interface of wildlife detection model after being prompted with an input image.




Figure 13. Interface of the Wiki page. The facts are being ordered by the confidence of each detection.

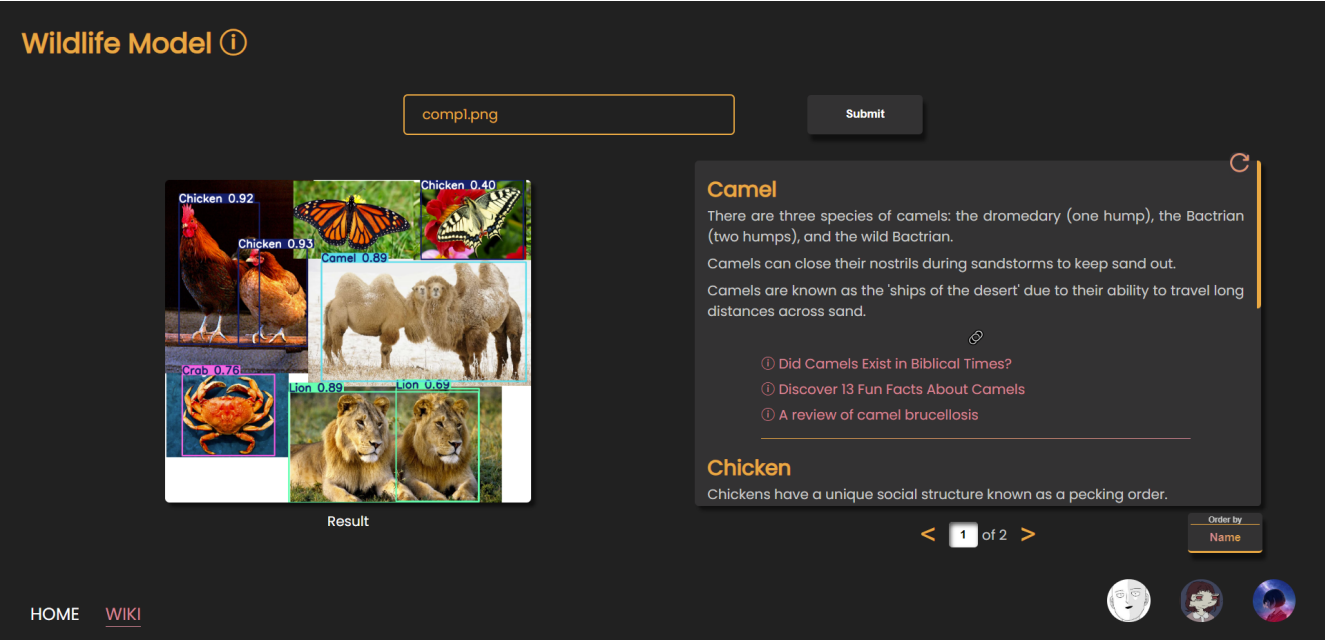


Figure 14. Same as Figure 13, but being ordered alphabetically by animal names.

ACKNOWLEDGMENT

The authors would like to thank the lecturers of the *Statistical Machine Learning* course at the Ho Chi Minh University of Science (Le Long Quoc, and Ngo Minh Nhut) for their instructions during theory and lab sessions as well as support during the project's duration.

The structure and style of this report is based on the Monthly Notices of the Royal Astronomical Society (MNRAS) template and instructions for authors (<https://academic.oup.com/mnras/>). All data and codes used in this project are for educational and research purposes.

REFERENCES

- Bochkovskiy A., Wang C.-Y., Liao H.-Y. M., 2020, YOLOv4: Optimal Speed and Accuracy of Object Detection ([arXiv:2004.10934](https://arxiv.org/abs/2004.10934))
- He K., Zhang X., Ren S., Sun J., 2015, [IEEE Transactions on Pattern Analysis and Machine Intelligence](https://doi.org/10.1109/TPAMI.2015.2462214), 37, 1904
- Liu S., Qi L., Qin H., Shi J., Jia J., 2018, Path Aggregation Network for Instance Segmentation ([arXiv:1803.01534](https://arxiv.org/abs/1803.01534))
- Redmon J., Divvala S., Girshick R., Farhadi A., 2016, in Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR).
- Tan M., 2019, EfficientNet: Rethinking model scaling for convolutional neural networks ([arXiv:1905.11946](https://arxiv.org/abs/1905.11946))
- Tianxiaomo 2021, pytorch-YOLOv4, <https://github.com/Tianxiaomo/pytorch-YOLOv4>
- Wang C.-Y., Liao H.-Y. M., Wu Y.-H., Chen P.-Y., Hsieh J.-W., Yeh I.-H., 2020, in Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops.